

Accelerating Optimal Power Flow with GPUs: SIMD Abstraction of Nonlinear Programs and Condensed-Space Interior-Point Methods

Sungho Shin, François Pacaud, and Mihai Anitescu

Abstract—This paper introduces a novel computational framework for solving alternating current (AC) optimal power flow (OPF) problems using graphics processing units (GPUs). While GPUs have demonstrated remarkable performance in various computing domains, their application in AC OPF has been limited due to challenges associated with porting sparse automatic differentiation (AD) and sparse linear solver routines to GPUs. We aim to address these issues with two key strategies. First, we utilize a single-instruction, multiple-data (SIMD) abstraction of nonlinear programs (NLP). This approach enables the specification of model equations while preserving their parallelizable structure, and in turn, facilitates the implementation of AD routines that can exploit such structure. Second, we employ a condensed-space interior-point method (IPM) with an inequality relaxation strategy. This technique involves relaxing equality constraints to inequalities and condensing the Karush-Kuhn-Tucker system into a much smaller positive definite matrix. This strategy offers the key advantage of being able to factorize the KKT matrix without numerical pivoting, which in the past has hampered the parallelization of the algorithm. By combining these two strategies, we can perform the majority of operations on GPUs while keeping the data residing in the device memory only. Comprehensive numerical benchmark results showcase the substantial computational advantage of our approach. Remarkably, for solving large-scale AC OPF problems to a moderate accuracy, our implementations—MadNLP.jl and SIMDiff.jl—running on NVIDIA GPUs achieve an order of magnitude speedup compared to state-of-the-art tools running on contemporary CPUs.

Index Terms—nonlinear programming, automatic differentiation, GPU computing, optimal power flow

I. INTRODUCTION

While graphics processing units (GPUs) have demonstrated remarkable capabilities in various computing domains, their adoption in large-scale constrained nonlinear programming (NLP) regimes, such as alternating current (AC) optimal power flow (OPF) problems, has faced some limitations. One of the primary challenges arises from the automatic differentiation (AD) of sparse model equations and the parallel factorization of indefinite sparse matrices, which are commonly encountered within constrained numerical optimization tasks [1]. While GPU computation can trivially accelerate several parts of the optimization process, especially various internal computations within the optimization solver, the sluggish data transfer between host and device memory hampers the ad-hoc implemen-




Fig. 1. A schematic description of host (CPU) and device (GPU) memory.

tation of GPU accelerations (Fig. II-C). To fully leverage the potential that modern GPU hardware has to offer, it becomes imperative to have a comprehensive computational framework for optimization on GPUs. That is, we need an AD/algebraic modeling framework, sparse linear solvers, and NLP solvers that can operate entirely on the GPU. Specifically, for the best performance, both the problem data and the solver's intermediate computational data must be exclusively resident within the device memory, with the majority of operations executed on the GPU.

This paper presents our approach to implement a comprehensive computational framework for solving large-scale AC OPF problems on NVIDIA GPUs, along with the associated software implementations: SIMDiff.jl [2], an algebraic modeling/AD tool, and MadNLP.jl [3], an NLP solver. Our approach incorporates two novel strategies: (i) a single-instruction, multiple-data (SIMD) abstraction of nonlinear programs (NLPs), enabling streamlined parallel AD on GPUs, and (ii) a condensed-space interior-point method (IPM) with an inequality relaxation strategy, which facilitates the use of highly efficient *refactorization* routines for sparse matrix factorization with fixed pivot sequences.

While derivative evaluation can be generally cheaper than linear algebra operations, our numerical results show that AD often constitutes more than half of the total solver time when using off-the-shelf AD implementations like JuMP.jl or AMPL. Instead of the general-purpose tools, our method leverages a specialized AD implementation based on the SIMD abstraction of NLPs. This abstraction allows to preserve the parallelizable structure within the model equations, facilitating

Sungho Shin and Mihai Anitescu are with Mathematics and Computer Science Division, Argonne National Laboratory, Lemont, IL, USA. e-mail: sshin@anl.gov, anitescu@mcs.anl.gov

François Pacaud is with Centre Automatique et Systèmes, Mines Paris - PSL, Paris, France. e-mail: francois.pacaud@minesparis.psl.eu

efficient derivative evaluations on the GPU. The AC power flow model is particularly well-suited for this abstraction as it involves repetitive expressions for each component type (e.g., buses, lines, generators), and the number of computational patterns does not increase with the network's size. Numerical results reported in this paper demonstrate that our proposed strategies can achieve over 20 times speedup by running on GPU. In comparison to general AD implementations on CPUs (such as AMPL and JuMP.jl), our GPU-based differentiation method can be approximately 50 times faster.

Linear algebra operations, especially sparse indefinite matrix factorization, are typically the bottleneck in NLP solution methods. Parallelizing this operation has been considered to be challenging, primarily due to the need for numerical pivoting, which is sequential in nature. However, when the matrix can be factorized without numerical pivoting, a significant part of the operation can be parallelized, and the numerical factorization can be efficiently performed on GPUs. We develop a condensed-space IPM strategy that allows the use of sparse matrix factorization routines without numerical pivoting. This strategy relaxes equality constraints by permitting small violations, which enables expressing the Karush-Kuhn-Tucker (KKT) system entirely in the primal space through the condensation procedure. Although this strategy is not new [4], it has traditionally been considered less efficient than the standard full-space method due to increased nonzero entries in the KKT system. However, when implemented on GPUs, it offers the key advantage of ensuring positive definiteness in the condensed KKT system. This, in turn, allows for the utilization of linear solvers with a fixed numerical pivot sequence (known as refactorization), an efficient implementation of which is available in the CUDA library. Although this method is susceptible to numerical stability issues due to increased condition number in the KKT system, our results demonstrate that the solver is robust enough to solve problems with a relative accuracy of $\epsilon_{\text{mach}}^{1/4} \approx 10^{-4}$.

We present numerical benchmark results to showcase the efficiency of our method, utilizing our implementations: MadNLP.jl and SIMDiff.jl. The solution of the KKT system is performed using the external cuSOLVER library. To assess the performance of our method, we compare it against standard CPU approaches using the data available in pglib-opf [5]. Our benchmark results demonstrate that our proposed computational framework has significant potential for accelerating the solution of AC OPF problems, especially when a moderate accuracy (e.g., 10^{-4}) is sufficient. Specifically, when running on NVIDIA GPUs, our method achieves a notable 4x speedup compared to our solver running on CPU, in the case of the largest instance. Moreover for the same instance, our approach surpasses the performance of existing tools (such as Ipopt interfaced with JuMP.jl) by an order of magnitude. This finding underscores the importance of harnessing the power of GPUs for tackling the computational challenges in power systems.

Related Work: Several recent works have explored the use of GPUs for large-scale nonlinear optimization problems. Cao et al. [6] proposed an augmented Lagrangian interior-point approach that employs augmented Lagrangian outer iteration and the treatment of linear systems using a preconditioned

conjugate gradient method. Another recent development is the hybrid (direct-iterative) KKT system solver specifically designed for GPUs [7]. Before the introduction of the condensed space IPM with an inequality relaxation strategy, the authors have investigated the use of reduction strategies (state variable elimination) to treat KKT systems in a dense form on the GPU [8]–[11]. An NLP solver for high-performance computers (HPC) with GPU accelerators called HiOP has been under development [12], which has a similar scope as our solver MadNLP.jl.

Organization: The paper is organized as follows. In the remainder of the current section, we introduce the mathematical notation. In Section II, we provide general preliminary knowledge on numerical optimization and GPU computing. In Section III, we present the SIMD abstraction of NLPs and their advantages in terms of implementing parallel AD. Section IV presents the optimization algorithm under study, the condensed-space IPM with an inequality relaxation strategy. Section V presents the numerical results, comparing our approach with other state-of-the-art solution methods on the CPUs. Finally, conclusions and future outlooks are given in Section VI.

Notation: We denote the set of real numbers and the set of integers by \mathbb{R} and \mathbb{I} . We let $[M] := \{1, 2, \dots, M\}$. We let $[v_i]_{i \in [M]} := [v_1; v_2; \dots, v_M]$. A vector of ones with an appropriate size is denoted by $\mathbf{1}$. An identity matrix with an appropriate size is denoted by I . For matrices A and B , $A \succ (\succeq) B$ indicates that $A - B$ is positive (semi)-definite while $A > (\geq) B$ denotes a componentwise inequality. We use the following convention for any symbol x : $X := \text{diag}(x)$.

II. PRELIMINARIES

This section covers three essential background topics: numerical optimization, AD, and GPU computing.

A. Numerical Optimization

We consider NLPs of the following form:

$$\min_{x^b \leq x \leq x^u} f(x) \quad (1a)$$

$$\text{s.t. } g(x) = 0. \quad (1b)$$

Numerous solution algorithms have been developed in the NLP literature. In terms of strategies to deal with inequality constraints, the NLP solution algorithms can be broadly classified into active-set methods and IPMs [4]. Active-set methods aim to find the set of active constraints associated with the optimal solution in a combinatorial manner, while IPMs replace inequality constraints with smooth barrier functions, enabling the use of Newton-type second-order algorithms. IPMs are known to be more scalable for problems with a large number of constraints and suitable for parallelization, thanks to the fixed sparsity pattern of the KKT matrix. Given these advantages, we have chosen IPMs as the backbone algorithm for developing our optimization methods on GPUs.

In terms of practical computations, three key components play vital roles: derivative evaluations (often provided by the AD capabilities of the algebraic modeling languages), linear

algebra operations, and various internal computations within the solver. Notably, most of the computational efforts are delegated to the external linear solver and AD library, while the optimization solver orchestrates the operation of these tools to drive the solution iterate towards the stationary point of the optimization problem.

Since the successful implementation of the open-source IPM solver in Ipopt, many subsequent implementations of NLP solvers [13]–[15] have based their implementation on Ipopt [5]. We also use Ipopt as our main reference for the IPM implementation. Below, we outline the overall computational procedure employed within the NLP solution frameworks:

1) Given the current primal-dual iterate $(x^{(\ell)}, y^{(\ell)}, z^b(\ell), z^\#(\ell))$, the AD package evaluates:

$$\nabla_x f(x^{(\ell)}), \quad \nabla_x g(x^{(\ell)}), \quad \nabla_{xx}^2 \mathcal{L}(x^{(\ell)}, y^{(\ell)}, z^b(\ell), z^\#(\ell)),$$

where

$$\mathcal{L}(x, y, z^b, z^\#) := f(x) - y^\top g(x) - z^b(x - x^b) - z^\#(x^\# - x).$$

2) The following sparse indefinite system (known as the KKT system) is solved using sparse indefinite factorization (typically, via sparse LBL^T factorization) and triangular solve routines.

$$\begin{bmatrix} W^{(\ell)} + \Sigma^{(\ell)} + \delta_w^{(\ell)} I & A^{(\ell)\top} \\ A^{(\ell)} & -\delta_c^{(\ell)} I \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} = \begin{bmatrix} r_x^{(\ell)} \\ r_y^{(\ell)} \end{bmatrix}, \quad (2)$$

where

$$\begin{aligned} W^{(\ell)} &:= \nabla_{xx}^2 \mathcal{L}(x^{(\ell)}, y^{(\ell)}, z^b(\ell), z^\#(\ell)), \quad A^{(\ell)} := \nabla_x g(x^{(\ell)}) \\ \Sigma^{(\ell)} &:= (X^{(\ell)})^{-1} Z^{(\ell)} \\ r_x^{(\ell)} &:= \nabla_x f(x^{(\ell)}) - \mu (X^{(\ell)})^{-1} \mathbf{1}, \quad r_y^{(\ell)} := g(x^{(\ell)}), \end{aligned}$$

and $\delta_w^{(\ell)}, \delta_c^{(\ell)} > 0$ are the regularization parameters determined based on the inertia correction procedure.

3) The optimization solver employs the filter line search procedure to determine the step size. The primal-dual iterate is updated by applying the determined step size and direction. This process is repeated until convergence criteria (typically based on the residual to the first-order optimality conditions) are satisfied.

B. AD

Numerical differentiation of computer programs can be achieved through three different methods: (i) finite difference method, (ii) symbolic differentiation, and (iii) AD. The finite difference method suffers from numerical rounding errors, and its computational complexity grows unfavorably with respect to the number of function arguments, making it less preferable unless no other alternatives are available. Symbolic differentiation uses computer algebra systems to obtain symbolic expressions of first or higher-order derivatives. While this method can differentiate functions up to high numerical precision, it suffers from "expression swelling" effect and struggles to compute the derivatives of long nested expressions in a computationally efficient way.

In contrast, AD differentiates computer programs by leveraging the computation graph. By inspecting the computation

graph and applying chain rules, AD can efficiently and accurately evaluate derivatives. This approach has become the dominant paradigm for derivative computation within the scientific computing domain, including NLP and machine learning. For large-scale optimization, such as AC OPFs, AD tools are often implemented as part of domain-specific modeling languages, such as JuMP, CasADi, and AMPL (optimization) and Torch and Flux (machine learning).

There are two alternative ways of propagating derivatives through the recursive application of chain rules: (i) forward-mode and (ii) reverse-mode, which operate in opposite directions (from leaves to root and from root to leaves, respectively). Reverse-mode automatic differentiation, also known as an adjoint method, has proven to be particularly effective for dealing with the function expressions in large-scale optimization problems.

The Julia Language, our language of choice, offers convenient and efficient ways to implement automatic differentiation. Through the use of the multiple dispatch paradigm, *any Julia function*, including commonly used operations like addition, multiplication, trigonometric and exponential functions, among others, can be easily overloaded. Multiple dispatch allows functions to be dynamically dispatched based on the run-time type, a crucial feature for implementing differentiable programming. Several AD implementations have been developed in Julia Language, such as ReverseDiff.jl, ForwardDiff.jl, Zygote.jl, and JuMP.jl. While these tools are general and useful for various applications, they are not optimized for evaluating derivatives of AC OPF problems, as they are not designed to exploit the parallelizable structures in the model.

C. GPU Computing

With the increasing prevalence of GPU in various scientific computing domains, there has been growing interest in leveraging these emerging architectures to efficiently solve large-scale NLPs, like AC OPF problems. However, adapting the NLP solution algorithms, such as IPM, to GPUs presents challenges due to the fundamental differences between GPU and CPU programming paradigms. While CPUs execute a sequence of instructions on a single input (single instruction, single data, or SISD, in Flynn's taxonomy), GPUs run the same instruction simultaneously on hundreds of threads using the SIMD paradigm (see Fig. II-C). The SIMD parallelism works well for algorithms that can be decomposed into simple instructions running entirely in parallel, but not all algorithms fit this paradigm. For example, branching in the control flow can hinder lockstep execution across multiple threads, and in turn, prevent efficient implementations on GPUs. On the other hand, when the algorithm's structure allows for efficient parallelization, the SIMD parallelism in GPUs can offer orders of magnitudes speedup.

We highlight that the following, arguably quite common, computational patterns are particularly effective when implemented on GPUs:

$$y \leftarrow [\phi(x; q_j)]_{j \in [J]} \quad (\text{Pattern 1})$$

$$o \leftarrow \text{Op } \psi(x; p_i)_{i \in [I]} \quad (\text{Pattern 2})$$

$$x \leftarrow \chi_{s_1} \circ \cdots \circ \chi_{s_K}(x) \quad (\text{Pattern 3})$$

Here, $\psi : \mathbb{R}^{n_x} \times \mathbb{R}^{n_p} \rightarrow \mathbb{R}$, $\phi : \mathbb{R}^{n_x} \times \mathbb{R}^{n_q} \rightarrow \mathbb{R}$, and $\chi : \mathbb{R}^{n_x} \times \mathbb{R}^{n_s} \rightarrow \mathbb{R}^{n_x}$ are simple instructions that require only a small number of operations; Op is a monoid operator on $\mathbb{R} \cup \{+\infty, -\infty\}$, such as addition, multiplication, maximum, and minimum; and the notation $\chi_{s_k}(x) := \chi(x, s_k)$ is used, and assume that \circ is commutative on $\{\chi_{s_k}(\cdot)\}_{\forall s_k}$. **Pattern 1** is typically most effective on GPUs, where each thread employed can operate independently without needing to simultaneously manipulate the same device memory location. While **Pattern 2** and **Pattern 3** is less effective as they require such a manipulation, but they still can be significantly faster than operations in CPUs, as substantial part of the operation can still be parallelized by the use of buffers. In simple cases, the implementation of these operations can be performed with the standard `map` and `mapreduce` programming models. However, in more complex cases, especially for **Pattern 3**, the implementation may require pre-inspection of memory write-access patterns and the use of custom kernels.

Many of the operations required in AD of sparse physical models, as well as the application of optimization algorithms, are based on the computational patterns mentioned above. For example, an AC OPF model can be implemented with 15 different computational patterns, most of which fall within the aforementioned categories. Furthermore, the computation within optimization solvers, such as forming the left-hand-side for the KKT systems, computing the $\|\cdot\|_\infty$ norm of the constraint violation, and assembling the condensed KKT system, can be carried out by using these computational patterns as well. The only exception is the factorization of the sparse KKT matrix, which requires more sophisticated implementations.

Implementing kernel functions for the above computational patterns in the Julia Language is straightforward, as it provides excellent high-level interfaces for array and kernel programming for GPU arrays. The code can even be device-agnostic, thanks to the portable programming capabilities brought by `KernelAbstractions.jl`. All of the AD and optimization capabilities in our tools, `MadNLP.jl` and `SIMDiff.jl`, are implemented in Julia Language, by leveraging its kernel and array programming capabilities.

III. SIMD ABSTRACTION OF NLPs

This section describes our implementation of SIMD abstraction and sparse AD of the model equations. The abstraction and AD are implemented as part of our algebraic modeling language `SIMDiff.jl`.

A. Abstraction

The SIMD abstraction under consideration is as follows:

$$\begin{aligned} \min_{x^b \leq x \leq x^\#} \sum_{l \in [L]} \sum_{i \in [I_l]} f^{(l)}(x; p_i^{(l)}) \quad (3) \\ \text{s.t. } \forall m \in [M] : \\ \left[g^{(m)}(x; q_j) \right]_{j \in [J_m]} + \sum_{n \in [N_m]} \sum_{k \in [K_n]} h^{(n)}(x; s_k^{(n)}) = 0, \end{aligned}$$




Fig. 2. A schematic description of SIMD parallelism

where $f^{(\ell)}(\cdot, \cdot)$, $g^{(m)}(\cdot, \cdot)$, and $h^{(n)}(\cdot, \cdot)$ are twice differentiable functions with respect to the first argument, $\{\{p_i^{(k)}\}_{i \in [N_k]}\}_{k \in [K]}$, $\{\{q_i^{(k)}\}_{i \in [M_i]}\}_{m \in [M]}$, and $\{\{s_k^{(n)}\}_{k \in [K_n]}\}_{n \in [N_m]}\}_{m \in [M]}$ are problem data, which can either be discrete or continuous. We assume that our functions $f^{(\ell)}(\cdot, \cdot)$, $g^{(m)}(\cdot, \cdot)$, and $h^{(n)}(\cdot, \cdot)$ can be expressed in the form of computational graphs of moderate length. One can observe that the problem in (3) is expressed by the computational patterns in Section II-C. In particular, the objective function falls within **Pattern 1**, the first term in the constraint falls within **Pattern 2**, and the second term in the constraint falls within **Pattern 3**. Accordingly, the evaluation and differentiation of the model equations in (3) is amenable to SIMD parallelism.

To implement the SIMD abstraction in the modeling environment, the algebraic modeling interface in `SIMDiff.jl` enforces the users to specify the model equations in an `Iterable` data type in Julia Language. This composite data type consists of an instruction (a Julia function) and data (a host or device array) over which the instruction is executed. This naturally facilitates maintaining the NLP model information in the form of SIMD abstraction in (3), and facilitates the model evaluation and differentiation on GPU accelerators.

B. Parallel AD

Many of the physics-based models, such as AC OPF, have highly repetitive structure. One of the manifestations of it is that the mathematical statement of the model is concise, even if the practical model may contain millions of variables and constraints. This is possible due to the use of repetition over a certain index and data sets. For example, it suffices to use 15 computational patterns to fully specify the AC OPF model. These patterns arise from (1) generation cost, (2) reference bus voltage angle constraint, (3-6) active and reactive power flow (to and from), (7) voltage angle difference constraint, (8-9) apparent power flow limits (to and from), (10-11) power balance equations, (12-13) generators' contributions to the power balance equations, and (14-15) in/out flows

contributions to the power balance equations. However, such repetitive structure is not well exploited during AD in the standard NLP modeling paradigms. In fact, without the SIMD abstraction, it is difficult for the AD package to detect the parallelizable structure within the model, as it will require the full inspection of the computational graph. By preserving the repetitive structures in the model, one can facilitate the implementation of parallelized derivative computations.

Using the multiple dispatch feature of Julia, `SIMDiff.jl` generates highly efficient derivative computation code, specifically compiled for each computational pattern in the model. These derivative evaluation code can be run over the data in various GPU array formats, implemented via array and kernel programming in Julia Language. In turn, `SIMDiff.jl` has the capability to efficiently evaluate first and second order derivatives using GPU accelerators.

C. Sparsity Analysis

The sparsity analysis is needed for determining the sparsity pattern of the evaluated derivatives. In the case of large-scale sparse problems, the initial sparsity analysis of nonlinear expressions can be expensive, as the sparsity should be analyzed for potentially millions of objective terms and constraints. However, oftentimes these analyses are applied for the same computational patterns, and the time and memory spent for sparsity analysis can be significantly reduced if the repetitive structures are exploited.

`SIMDiff.jl` exploits the SIMD abstraction of the model equations to save the computational cost spent for sparsity analysis. This is accomplished by applying sparsity analysis for the instruction for each computational pattern and expand the obtained sparsity pattern over the data over which the instruction is executed. Specifically, the sparsity analysis code exploits Julia's multiple dispatch feature to obtain parameterized sparsity pattern for each instruction, and the obtained parameterized sparsity pattern is materialized once the data array is given. This process drastically save the computational cost for the sparsity analysis.

IV. CONDENSED-SPACE IPMS WITH AN INEQUALITY RELAXATION STRATEGY

We present the condensed-space IPM within the context of the general NLP formulation in (1). Our method has two key differences from standard IPM implementations: (i) the use of inequality relaxation and (ii) the condensed treatment of the KKT system.

A. Inequality Relaxation

At the beginning of the algorithm, we apply inequality relaxation to replace the equality constraints in (1) with inequalities by introducing slack variables $s \in \mathbb{R}^m$:

$$g(x) - s = 0, \quad s^b \leq s \leq s^\sharp, \quad (4)$$

where $s^b, s^\sharp \in \mathbb{R}^m$ are chosen to be close to zero. This relaxed problem can be stated as follows:

$$\begin{aligned} \min_{\substack{x \\ [x^b] \leq [x] \leq [x^\sharp]}} f(x) \\ \text{s.t. } g(x) - s = 0. \end{aligned} \quad (5)$$

In our implementation, we set s^b, s^\sharp as $-\epsilon_{\text{tol}}\mathbf{1}$ and $+\epsilon_{\text{tol}}\mathbf{1}$, respectively, where $\epsilon_{\text{tol}} > 0$ is a user-specified relative tolerance of the IPM. This type of relaxation is commonly used in practical IPM implementations; for example, in `Ipopt`, the solver relaxes the bounds and inequality constraints by $O(\epsilon_{\text{tol}})$ to prevent an empty interior of the feasible set (see [5, Section 3.5]). For condensed-space IPM, we cannot maintain the same level of precision due to the increased condition number of the KKT system. We have found that setting ϵ_{tol} to be $\epsilon_{\text{mach}}^{1/4} \approx 10^{-4}$ ensures numerical stability while achieving satisfactory convergence behavior. Thus, our solver sets the tolerance to 10^{-4} by default when using condensed IPM.

B. Barrier Subproblem

The IPM replaces the equality and inequality-constrained problem in (4) with an equality-constrained barrier subproblem:

$$\min_x f(x) - \mu\mathbf{1}^\top \log(x - x^b) - \mu\mathbf{1}^\top \log(x^\sharp - x) \quad (6a)$$

$$- \mu\mathbf{1}^\top \log(s - s^b) - \mu\mathbf{1}^\top \log(s^\sharp - s) \quad (6b)$$

$$\text{s.t. } g(x) = 0.$$

Here, $\mu > 0$ is the barrier parameter. The smooth log-barrier function is employed to avoid handling inequalities in a combinatorial fashion (as in active set methods). A superlinear local convergence to the first-order stationary point can be achieved by repeatedly applying Newton's step to the KKT conditions of (6) with $\mu \searrow 0$.

C. Newton's Step Computation

The Newton step direction is computed by solving a so-called KKT system; to explain this, we consider the first-order optimality conditions (KKT conditions) for the barrier subproblem in (6):

$$\begin{aligned} \nabla_x f(x) - \nabla_x g(x)^\top y - z_x^b + z_x^\sharp &= 0 \\ -z_s^b + z_s^\sharp &= 0 & g(x) &= 0 \\ Z_x^b(x - x^b) - \mu\mathbf{1} &= 0 & Z_x^\sharp(x^\sharp - x) - \mu\mathbf{1} &= 0 \\ Z_s^b(s - s^b) - \mu\mathbf{1} &= 0, & Z_s^\sharp(s^\sharp - s) - \mu\mathbf{1} &= 0, \end{aligned} \quad (7)$$

where $y \in \mathbb{R}^m$, $z_x^b, z_x^\sharp \in \mathbb{R}^n$, and $z_s^b, z_s^\sharp \in \mathbb{R}^m$ are Lagrange multipliers associated with the equality and bound constraints in (5). The Newton step for solving the nonlinear equations in (7) can be computed by solving the KKT system in (8). Here, we recall the definitions of $W^{(\ell)}$ and $A^{(\ell)}$ from Section II-A, and $p_x^{(\ell)}, \dots, p_{z_s^\sharp}^{(\ell)}$ are defined by the left-hand-sides of the equations in (7). In sequel, we shall drop the superscript $(\cdot)^{(\ell)}$ for concise notation.

Now, we observe that a significant portion of the system in (8) can be eliminated by exploiting the block structure, leading to an equivalent system stated in a smaller space. In particular, the lower-right 4×4 block is always invertible since the IPM procedure ensures that the iterates stay in the strict interior

$$\underbrace{\begin{bmatrix} W^{(\ell)} + \delta_w^{(\ell)} I & A^{(\ell)\top} & -I & I & & \\ & \delta_w^{(\ell)} I & -I & & -I & I \\ & A^{(\ell)} & -I & -\delta_c^{(\ell)} I & & \\ Z_x^{(\ell)b} & & & X^{(\ell)} - X^b & & \\ -Z_x^{(\ell)\sharp} & & & & X^\sharp - X^{(\ell)} & \\ & Z_s^{(\ell)b} & & & & S^{(\ell)} - S^b \\ & -Z_s^{(\ell)\sharp} & & & & S^\sharp - S^{(\ell)} \end{bmatrix}}_{M_{\text{full}}} \begin{bmatrix} \Delta x \\ \Delta s \\ \Delta y \\ \Delta z_x^b \\ \Delta z_x^\sharp \\ \Delta z_s^b \\ \Delta z_s^\sharp \end{bmatrix} = \begin{bmatrix} p_x^{(\ell)} \\ p_s^{(\ell)} \\ p_y^{(\ell)} \\ p_{z_x^b}^{(\ell)} \\ p_{z_x^\sharp}^{(\ell)} \\ p_{z_s^b}^{(\ell)} \\ p_{z_s^\sharp}^{(\ell)} \end{bmatrix} \quad (8)$$

of the feasible set. This allows for eliminating the lower-right 4x4 block, resulting in:

$$\underbrace{\begin{bmatrix} W + \Sigma_x + \delta_w I & A^\top & -I \\ A & -I & -\delta_c I \end{bmatrix}}_{M_{\text{red}}} \begin{bmatrix} \Delta x \\ \Delta s \\ \Delta y \end{bmatrix} = \begin{bmatrix} q_x \\ q_s \\ q_y \end{bmatrix}, \quad (9)$$

where the dependence on the evaluation point is suppressed for concise notation, and:

$$\begin{aligned} \Sigma_x &:= Z_x^b (X - X^b)^{-1} + Z_x^\sharp (X^\sharp - X)^{-1} \\ \Sigma_s &:= Z_s^b (S - S^b)^{-1} + Z_s^\sharp (S^\sharp - S)^{-1} \\ q_x &:= p_x + (X - X^b)^{-1} p_{z_x^b} - (X^\sharp - X)^{-1} p_{z_x^\sharp} \\ q_s &:= (S - S^b)^{-1} p_{z_s^b} - (S^\sharp - S)^{-1} p_{z_s^\sharp} \\ q_y &:= p_y. \end{aligned}$$

The bound dual steps can be recovered as follows:

$$\begin{aligned} \Delta z_x^b &= (X - X^b)^{-1} (-Z_x^b \Delta x + p_{z_x^b}) \\ \Delta z_x^\sharp &= (X^\sharp - X)^{-1} (Z_x^\sharp \Delta x + p_{z_x^\sharp}) \\ \Delta z_s^b &= (S - S^b)^{-1} (-Z_s^b \Delta s + p_{z_s^b}) \\ \Delta z_s^\sharp &= (S^\sharp - S)^{-1} (Z_s^\sharp \Delta s + p_{z_s^\sharp}). \end{aligned} \quad (10)$$

Note that the matrices involved in the inversions in (10) are always diagonal, so their computation is not expensive. Also, note that the reduced system in (9) corresponds to the KKT system in (2). However, in the original version of the algorithm, we did not introduce the slack variables, so it does not have the additional structure imposed by the slack variables.

The key advantage of the inequality relaxation strategy is that it imposes additional structure to the reduced KKT system, allowing us to further reduce the dimension of the problem. In particular, the lower-right 2x2 block in (9) can be eliminated, which is a procedure called *condensation*; here, the invertibility of the lower-right block can be verified from the fact that $\delta_w, \delta_c \geq 0$ and $\Sigma_s \succ 0$. Through this, we obtain the following system, written in the primal space only:

$$\underbrace{(W + \delta_w I + \Sigma_x + A^\top D A)}_{M_{\text{cond}}} \Delta x = q_x + A^\top (C q_y + D q_s), \quad (11)$$

where:

$$C := (\delta_c \Sigma_s + (1 + \delta_c \delta_w) I)^{-1}, \quad D := (\Sigma_s + \delta_w I) C,$$

and the dual and slack step directions can be recovered by:

$$\begin{aligned} \Delta s &:= C (\delta_c q_s - (q_y + A \Delta x)) \\ \Delta y &:= (\Sigma_s + \delta_w I) \Delta s - q_s. \end{aligned} \quad (12)$$

Again, the matrices involved in the inversions above are always diagonal, so their computation is not expensive.

Therefore, the only sparse matrix that needs to be factorized is the matrix in (11), which is in $\mathbb{R}^{n \times n}$. In general NLPs, the condensation strategy can arbitrarily increase the density of the KKT system, but for AC OPF problems, the condensed KKT system is still sparse. This is because the maximum number of nonzeros per row in constraint Jacobian is bounded by the graph degree, which should be reasonably small in practice.

The reason that the condensation strategy is particularly relevant for GPUs is that the matrix in (11) is positive definite upon the application of standard inertia correction method. Typically, to guarantee that the computed step direction is a descent direction, we need a condition that $\text{inertia}(M_{\text{red}}) = (n + 5m, 0, m)$. Here, inertia refers to the tuple of positive, zero, and negative eigenvalues. Accordingly, we employ inertia correction methods to modify the augmented KKT system so that the desired conditions on the inertia is satisfied.

By Sylvester's law, we have:

$$\begin{aligned} \text{inertia}(M_{\text{red}}) &= (n + 5m, 0, m) \\ \iff \text{inertia}(M_{\text{cond}}) &= (n, 0, 0). \end{aligned}$$

Thus, any choice of $\delta_w, \delta_c > 0$ that makes the condensed KKT system positive definite yields the the desired inertia condition. An important observation here is that the condensed KKT matrix with desired inertia condition is always positive definite. Thus, M_{cond} can be factorized with fixed pivoting (e.g., Cholesky factorization or LU refactorization), which is significantly more amenable to parallel implementation than indefinite LBL^T factorization, which is standard in the state-of-the-art CPU algorithms but requires the use of pivoting.

Although we call (11) a condensed KKT system, M_{cond} is not necessarily a dense matrix. In fact, in the case of AC OPF problems, M_{cond} is still highly sparse, as W and A are graph-induced banded systems. Thus, exploiting sparsity is still necessary to enable scalable computations.

Algorithm 1 Condensed-Space IPM

Require: Primal-dual solution guesses x, y, z^b, z^\sharp , bounds $x^b, x^\sharp, s^b, s^\sharp$, callbacks $f(\cdot), g(\cdot), \nabla_x f(\cdot), \nabla_x g(\cdot), \nabla_{xx}^2 \mathcal{L}(\cdot)$, and tolerance ϵ_{tol}

- 1: Relax the equality constraints by (4) and initialize the slack s and the associated dual variables z_s^b, z_s^\sharp .
- 2: **while** convergence criteria (14) not satisfied **do**
- 3: Solve the condensed KKT system (11) with $\delta_w = \delta_c = 0$ to compute the primal step Δx and recover the dual steps $\Delta y, \Delta z_x^b, \Delta z_x^\sharp, \Delta z_s^b, \Delta z_s^\sharp$ by (10) and (12).
- 4: Determine the need for regularization and if necessary, recompute the step directions with proper choices of $\delta_w, \delta_c > 0$.
- 5: Choose step sizes $\alpha, \alpha_z > 0$ via line search.
- 6: Update the solution by (13).
- 7: Update filter and barrier parameter μ .
- 8: **end while**
- 9: **return** The first-order stationary points $x^*, y^*, z^{b*}, z^{\sharp*}$

Due to the fact that the solutions of NLPs often have multiple active constraints, and the presence of active constraints can arbitrarily increase the conditioning of the KKT System, the system in (11) can be severely ill-conditioned. As such, a single triangular solve may not provide a sufficiently accurate step direction. Accordingly, iterative refinement methods are employed to refine the solution by performing multiple triangular solves. Notably, iterative refinement is applied to the full KKT system, rather than the condensed system in (11).

D. Line Search and IPM Iterations

The step size can be determined using the line search procedure. Although there are numerous alternative approaches, we follow the filter line search method implemented in the Ipopt solver. The line search procedure employed here determines the step size by performing backtracking line search until a trial point satisfying sufficient progress conditions are satisfied and acceptable by the filter. Furthermore, to enhance the convergence behavior, various additional strategies are implemented, such as the second-order correction, restoration phase, and automatic scaling. For the details of the implementation of the filter line search and various additional strategies, the readers are referred to [5].

The step size and direction obtained above can be implemented as follows:

$$\begin{aligned} (x, s, y) &\leftarrow (x, s, y) + \alpha(\Delta x, \Delta s, \Delta y), \\ (z_x^b, z_x^\sharp, z_s^b, z_s^\sharp) &\leftarrow (z_x^b, z_x^\sharp, z_s^b, z_s^\sharp) + \alpha_z(\Delta z_x^b, \Delta z_x^\sharp, \Delta z_s^b, \Delta z_s^\sharp). \end{aligned} \quad (13)$$

The iteration in (13) is repeated until the convergence criterion is satisfied. The convergence criterion is defined as

$$\text{residual}(x^{(\ell)}, s^{(\ell)}, y^{(\ell)}, z_x^{(\ell)b}, z_x^{(\ell)\sharp}, z_s^{(\ell)b}, z_s^{(\ell)\sharp}) < \epsilon_{\text{tol}}, \quad (14)$$

where $\text{residual}(\cdot)$ is a scaled version of the residual to the first-order conditions in (7). Finally, we summarize our condensed-space IPM in Algorithm 1.

E. Notes on the Implementation

We have implemented the condensed-space IPM by adapting our code base in MadNLP.jl, which implements the IPM as in Ipopt in Julia Language. A key feature of MadNLP is that the IPM is implemented with high levels of abstractions, while the specific handling of the data structures within the KKT systems is carried out by data-type specific kernel functions. This design allows us to apply the mathematical operations equivalent to Ipopt to different data structures, such as SparseKKTSystem, DenseKKTSystem, DenseCondensedKKTSystem, etc, whose data are stored either on host or device memory. For the implementation of the condensed-space IPM presented in this paper, we have added a new type of KKT system called SparseCondensedKKTSystem and implemented additional kernels needed for handling the data structures specific to this KKT system type. This approach ensures that we are performing mathematically equivalent operations as in the mature, extensively tested existing code base. This also allows us to easily switch between different KKT system types, which is crucial for experimenting with various solvers and data structures, as well as for efficiently leveraging GPU acceleration when available. Furthermore, by maintaining this level of abstraction, the condensed-space IPM can be seamlessly integrated into the existing framework, making it easier to maintain and extend in the future.

V. NUMERICAL RESULTS

This section presents the numerical benchmark results, comparing our method against state-of-the-art methods on CPUs for solving standard AC OPF problems.

A. Methods

We compared four different configurations of NLP solution frameworks:

- MadNLP.jl + SIMDiff.jl + cuSOLVER (GPU) (Config 1)
- MadNLP.jl + SIMDiff.jl + Ma27 (CPU) (Config 2)
- Ipopt + AMPL + Ma27 (CPU) (Config 3)
- Ipopt + JuMP.jl + Ma27 (CPU) (Config 4)

Config 1 is our main GPU configuration; Config 2 represents our implementation running on CPU; and Config 3 and Config 4 are used as benchmarks. Config 1 and Config 2 share significant amount of code, especially the high-level abstractions, but they differ in how they handle the KKT systems. In Config 1, MadNLP.jl applies the condensed-space IPM along with the inequality relaxation strategy, while in Config 2, MadNLP.jl applies IPM based on the indefinite, non-condensed KKT system, as in (2). In Config 1, we use the cuSOLVER library to solve the condensed KKT system. The initial symbolic factorization is performed using the KLU package [16], and the subsequent numerical factorization and triangular solves are performed by cuSOLVER with the fixed pivot sequence from KLU. Software and hardware details of each configuration are illustrated in Table I. The AC OPF




Fig. 3. Speedup achieved by using GPUs.

problem is formulated using the model from the rosetta-opf project [17], and the test cases are obtained from the pglb-opf repository [18]. The external packages are called from Julia Language, through thin wrapper packages, such as Ipopt.jl and AmplNLWriters.jl. A tolerance of 10^{-4} is set for MadNLP.jl and Ipopt solvers, with other solver options adjusted to ensure a fair comparison across different solvers. The results can be reproduced with the script available at <https://github.com/sshin23/opf-on-gpu>.

B. Results

The numerical benchmark results, including total solution time and its breakdown into linear algebra and derivative evaluation time (with the remainder considered as solver internal time), are shown in Table II. The quality of the solution (objective value and constraint violation measured by $\|\cdot\|_\infty$) is shown in Table III. Figure 3 visually represents the speedup brought by GPUs, by comparing the timing results of Config 1 and Config 2.

First, when comparing the solvers' performance in terms of the interior point method (IPM) iteration counts, MadNLP.jl is as efficient as the state-of-the-art solver Ipopt. The IPM iteration count is nearly the same as that of Ipopt for achieving the same level of accuracy in the final solution (see Table III). This suggests that running mathematically equivalent operations on GPUs (by using shared code base in high-level abstractions) can yield a similar degree of effectiveness in terms of IPM convergence.

Next, we discuss the effectiveness of parallel AD on GPUs. We observe that even on CPUs, SIMDiff.jl is substantially faster than the AD routines implemented in AMPL or JuMP.jl. This efficiency arises from SIMDiff.jl obtaining the derivative functions specifically compiled for the type of model, which is optimized for the distinctive computational pattern in the model. When comparing SIMDiff.jl running on CPUs and GPUs, we observe a further speedup of up to almost 20x for large instances (e.g., case24465_goc); remarkably, this is 50 times faster than JuMP.jl. This demonstrates that parallelization of AD brings significant computational gain. This is because most operations in the derivative evaluations are Pattern 1 operations, which are highly effective on GPUs.

We next discuss the linear solver time. While the speedup achieved by linear solvers is only moderate, this has a high

impact on the overall speedup, as linear solver time constitutes a significant portion of the total solution time. For large instances, approximately 4x speedup can be achieved, though this can vary depending on the instances. For example, for case24464_goc, the GPU linear solver was slower than the CPUs. The investigation of under which circumstances cuSOLVER is more effective warrants further research.

Solver internal time could also be significantly accelerated through GPU utilization. We can observe that the speedup in solver internal operations is consistently greater than the speedup in linear solvers. However, due to the frequent use of Pattern 2 operations, the speedup in solver internal operations is lesser than that of derivative evaluations.

Overall, our GPU implementation exhibits significant speedup across all components: derivative evaluation, linear algebra, and solver internal computation, resulting in substantial gains in total solution time. The results indicate that GPUs become more effective for large-scale instances, particularly when the number of variables is greater than 20,000. Notably, for the largest instance, case30000_goc, our GPU implementation is 4 times faster than our CPU implementation and approximately 10 times faster than the solver configuration with state-of-the-art tools (Ipopt, JuMP.jl, and Ma27). This demonstrates that GPUs can bring significant computational gains for large-scale AC OPF problems, enabling the solution of previously inconceivable problems due to the limitations of existing CPU-based solution tools.

VI. CONCLUSIONS AND FUTURE OUTLOOK

We have presented an NLP solution framework for solving large-scale AC OPF problems. By leveraging the SIMD abstraction of NLPs and a condensed-space IPM, we have effectively eliminated the need for serial computations, enabling the implementation of a solution framework that can run entirely on GPUs. Our method has demonstrated promising results, achieving a 4x speedup when compared to CPU implementations for large-scale AC OPF problems. Notably, our approach outperforms one of the state-of-the-art CPU-based implementations by approximately 10 times. These results, along with our packages MadNLP.jl and SIMDiff.jl, showcase a significant advancement in our capabilities in dealing with large-scale optimization problems in power systems and underscore the potential of accelerated computing in large-scale optimization area. However, the condensation procedure leads to an increase in the condition number of the KKT system, resulting in decreased final solution accuracy. Addressing the challenges posed by ill-conditioning remains to be an important aspect for future work. In the following paragraphs, we discuss some remaining open questions and future outlooks.

Obtaining Higher Numerical Precisions: While we have focused on the IPM, other constrained optimization paradigms, such as penalty methods and augmented Lagrangian methods exist, and similar strategies based on condensed linear systems can be developed. It would be valuable to investigate which algorithm would be the right paradigm for constrained large-scale optimization on GPU that can best handle the ill-

TABLE I
DETAILS OF NUMERICAL EXPERIMENT SETTINGS

	MadNLP.jl + SIMDiff.jl + cuSOLVER (gpu)	MadNLP.jl + SIMDiff.jl + Ma27 (cpu)	Ipopt + AMPL + Ma27 (cpu)	Ipopt + JuMP.jl + Ma27 (cpu)
Optimization Solver	MadNLP.jl (dev)*		Ipopt (v3.13.3)	
Derivative Evaluations	SIMDiff.jl (dev)*		AMPL Solver Library	
Linear Solver	cuSOLVER (v11.4.5)		Ma27 (v2015.06.23)	
Hardware	NVIDIA Quadro GV100		Intel Xeon Gold 6140	

*Specific commit hashes are available at <https://github.com/sshin23/opf-on-gpu>

TABLE II
NUMERICAL RESULTS

Case	nvars	ncons	MadNLP + SIMDiff + cuSOLVER (gpu)				MadNLP + SIMDiff + Ma27 (cpu)				Ipopt + AMPL + Ma27 (cpu)				Ipopt + JuMP + Ma27 (cpu)			
			iter	deriv. [†]	lin. [‡]	total [†]	iter	deriv. [†]	lin. [‡]	total [†]	iter	deriv. [‡]	total [‡]		iter	deriv. [‡]	total [‡]	
89_pegase	1.0k	1.6k	28	0.03	0.23	0.33	31	0.01	0.03	0.06	29	0.04	0.08		29	0.11	0.17	
179_goc	1.5k	2.2k	30	0.04	0.61	0.74	43	0.01	0.05	0.09	42	0.05	0.11		42	0.15	0.24	
500_goc	4.3k	6.1k	36	0.04	0.45	0.58	35	0.02	0.12	0.20	36	0.13	0.30		34	0.41	0.61	
793_goc	5.4k	8.0k	33	0.03	0.36	0.49	31	0.02	0.15	0.24	31	0.19	0.37		30	0.61	0.84	
1354_pegase	11.2k	16.6k	47	0.08	1.07	1.35	45	0.07	0.42	0.70	41	0.91	1.43		41	2.36	3.02	
2312_goc	17.1k	25.7k	38	0.04	1.16	1.33	40	0.10	0.74	1.13	38	1.45	2.33		38	3.16	4.14	
2000_goc	19.0k	29.4k	36	0.04	0.99	1.18	38	0.11	0.82	1.29	39	1.73	2.76		38	4.19	5.32	
3022_goc	23.2k	35.0k	43	0.04	1.39	1.63	49	0.18	1.27	1.93	47	2.56	4.02		47	5.68	7.29	
2742_goc	24.5k	38.2k	155	0.26	4.54	5.54	122	0.62	5.31	7.60	97	8.22	13.66		98	20.09	26.02	
2869_pegase	25.1k	37.8k	52	0.05	1.70	1.97	52	0.21	1.56	2.35	50	3.19	4.89		50	6.07	8.00	
3970_goc	35.3k	54.4k	44	0.05	1.64	1.91	45	0.26	2.77	3.75	60	5.49	10.04		43	7.20	10.92	
4020_goc	36.7k	57.0k	68	0.07	2.94	3.35	59	0.36	5.66	7.01	55	5.43	11.87		55	10.72	17.54	
4917_goc	37.9k	56.9k	48	0.05	1.74	2.07	57	0.34	2.79	4.07	53	5.03	7.90		53	9.84	13.07	
4601_goc	38.8k	59.6k	71	0.07	2.46	2.87	66	0.41	4.37	5.92	69	6.92	12.66		68	12.82	18.74	
4837_goc	41.4k	64.0k	57	0.06	2.37	2.72	56	0.39	3.79	5.24	56	6.50	10.94		56	12.70	17.61	
4619_goc	42.5k	66.3k	54	0.06	2.59	2.97	46	0.32	4.54	5.78	48	5.49	11.02		46	10.04	15.48	
10000_goc	76.8k	112.4k	56	0.06	2.63	3.11	77	0.89	9.54	13.00	74	14.02	24.18		74	25.13	36.46	
8387_pegase	78.7k	118.7k	64	0.12	6.08	6.87	70	0.89	9.44	12.96	69	14.23	23.55		69	26.40	36.74	
9591_goc	83.6k	130.6k	69	0.11	6.84	7.70	65	0.92	17.20	20.82	64	14.96	35.70		62	28.71	49.75	
9241_pegase	85.6k	130.8k	60	0.10	4.35	5.15	63	0.89	10.34	13.91	61	14.09	24.33		61	25.98	37.19	
10480_goc	96.8k	150.9k	70	0.13	13.19	14.26	66	1.05	18.19	22.40	64	16.93	38.04		63	33.53	56.04	
13659_pegase	117.4k	170.6k	63	0.12	6.10	7.15	58	1.08	12.91	17.35	64	19.70	35.66		64	35.45	52.99	
19402_goc	179.6k	281.7k	79	0.17	21.47	23.28	70	2.25	51.82	61.06	70	36.50	95.34		70	68.12	127.29	
24464_goc	203.4k	313.6k	63	0.11	69.32	70.63	58	2.22	33.03	41.71	58	33.50	70.15		58	62.04	102.17	
30000_goc	208.6k	307.8k	162	0.33	18.42	22.05	136	5.68	80.01	100.25	180	101.98	249.81		126	135.11	209.45	

[†]Wall time (sec) measured by Julia. [‡]CPU time (sec) reported by Ipopt.

conditioning issue of the condensed KKT system, and in turn, achieve the highest degree of accuracy.

Security-Constrained, Multi-Period, Distribution OPFs: Though the proposed method has demonstrated significant computational advantages for transmission AC OPF problems, our results can also interpreted that efficient CPUs can still handle these problems reasonably well. We anticipate that there will be more substantial performance gains for larger-scale optimization problems, such as security-constrained and multi-period OPFs or joint optimization problems involving transmission, distribution, and gas network systems. We are interested in exploring Schur complement-based decomposition approaches, combined with the condensation-based strategy, similarly to [8], to demonstrate even greater scalability.

Alternative Linear Solvers: While cuSOLVERRF has been effective for solving the condensed KKT systems using LU factorization, Cholesky factorization holds promise for better performance due to lower computational complexity and the ability to reveal the inertia of the KKT system. In the future, we are interested in exploring other linear solvers options, such as CHOLMOD [19], Baspacho [20], and HyKKT [7].

Portability: Our implementation is currently only tested on NVIDIA GPUs, but our GPU implementation is largely based on array programming and KernelAbstractions.jl in Julia, which are in principle compatible with various GPU

architectures, including AMD, Intel, and Apple GPUs. By incorporating cross-architecture linear solvers, we envision supporting broader class of GPU accelerators in the future.

REFERENCES

- [1] M. Anitescu, K. Kim, Y. Kim, A. Maldonado, F. Pacaud, V. Rao, M. Schanen, S. Shin, and A. Subramanian, "Targeting Exascale with Julia on GPUs for multiperiod optimization with scenario constraints," *SIAG/OPT Views and News*, 2021. [Online]. Available: <http://wiki.siam.org/siag-op/images/siag-op/e8/ViewsAndNews-29-1.pdf>
- [2] S. Shin, "SIMDiff.jl," <https://github.com/sshin23/SIMDiff.jl>.
- [3] S. Shin and F. Pacaud, "MadNLP.jl," <https://github.com/MadNLP/MadNLP.jl>.
- [4] J. Nocedal and S. J. Wright, *Numerical optimization*. Springer, 2006.
- [5] A. Wächter and L. T. Biegler, "On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming," *Mathematical programming*, vol. 106, pp. 25–57, 2006.
- [6] Y. Cao, A. Seth, and C. D. Laird, "An augmented lagrangian interior-point approach for large-scale nlp problems on graphics processing units," *Computers & Chemical Engineering*, vol. 85, pp. 76–83, 2016.
- [7] S. Regev, N.-Y. Chiang, E. Darve, C. G. Petra, M. A. Saunders, K. Świrnydowicz, and S. Peleš, "Hykkt: a hybrid direct-iterative method for solving kkt linear systems," *Optimization Methods and Software*, vol. 38, no. 2, pp. 332–355, 2023.
- [8] F. Pacaud, M. Schanen, S. Shin, D. A. Maldonado, and M. Anitescu, "Parallel interior-point solver for block-structured nonlinear programs on simd/gpu architectures," *arXiv preprint arXiv:2301.04869*, 2023.
- [9] F. Pacaud, D. A. Maldonado, S. Shin, M. Schanen, and M. Anitescu, "A feasible reduced space method for real-time optimal power flow," *Electric Power Systems Research*, vol. 212, p. 108268, 2022.

TABLE III
SOLUTION QUALITY

Case	MadNLP + SIMDiff + cuSOLVER (gpu)		MadNLP + SIMDiff + Ma27 (cpu)		Ipopt + AMPL + Ma27 (cpu)		Ipopt + JuMP + Ma27 (cpu)	
	objective	constr. viol.	objective	constr. viol.	objective	constr. viol.	objective	constr. viol.
89_pegase	1.07023029e+05	1.69977362e-03	1.07277300e+05	1.69995406e-03	1.07273132e+05	1.69762454e-02	1.07273132e+05	1.69762454e-02
179_goc	7.54098231e+05	3.64045772e-03	7.54215279e+05	3.64095371e-03	7.54214091e+05	1.05727439e-02	7.54214091e+05	1.05727439e-02
500_goc	4.53056588e+05	1.16442922e-03	4.54894607e+05	1.16461929e-03	4.54894301e+05	1.16449188e-03	4.54894349e+05	1.16443248e-03
793_goc	2.59660004e+05	1.12495280e-03	2.60179408e+05	1.14373500e-03	2.60177953e+05	2.52890328e-02	2.60177960e+05	2.52825510e-02
1354_pegase	1.25574315e+06	4.18838427e-03	1.25874608e+06	4.18894441e-03	1.25873160e+06	2.91106529e-02	1.25873160e+06	2.91106529e-02
2312_goc	4.40492687e+05	1.95782217e-03	4.41301927e+05	1.98487972e-03	4.41301012e+05	2.86441953e-03	4.41301012e+05	2.86441953e-03
2000_goc	9.66186544e+05	1.07957382e-03	9.73392385e+05	1.07991565e-03	9.73392524e+05	1.07970410e-03	9.73392602e+05	1.07958552e-03
3022_goc	6.00461469e+05	1.60590210e-03	6.01341340e+05	1.92264271e-03	6.01340934e+05	7.06720510e-03	6.01340934e+05	7.06720510e-03
2742_goc	2.70328757e+05	9.99725733e-04	2.75672815e+05	9.99997332e-04	2.75672759e+05	1.13868333e-03	2.75672759e+05	1.13868333e-03
2869_pegase	2.45584120e+06	4.18833905e-03	2.46259584e+06	4.18882610e-03	2.46258759e+06	3.15283321e-02	2.46258759e+06	3.15283321e-02
3970_goc	9.27998953e+05	6.41922608e-04	9.60666837e+05	6.42469892e-04	9.60667021e+05	6.42371530e-04	9.60667776e+05	6.41960999e-04
4020_goc	8.02565861e+05	1.29969745e-03	8.21952202e+05	1.29998686e-03	8.21952543e+05	1.29986624e-03	8.21952543e+05	1.29986624e-03
4917_goc	1.38537252e+06	1.54172485e-03	1.38769645e+06	1.70860688e-03	1.38769342e+06	1.62739725e-02	1.38769342e+06	1.62739725e-02
4601_goc	7.92510931e+05	9.99886244e-04	8.25898288e+05	9.99978318e-04	8.25898470e+05	9.99896654e-04	8.25898481e+05	9.99894295e-04
4837_goc	8.60071647e+05	9.92673673e-04	8.72192598e+05	9.92934504e-04	8.72192733e+05	9.92677263e-04	8.72192733e+05	9.92677263e-04
4619_goc	4.66738422e+05	8.80364611e-04	4.76659294e+05	8.80485073e-04	4.76659432e+05	8.80367536e-04	4.76659432e+05	8.80367536e-04
10000_goc	1.34739992e+06	5.36209615e-04	1.35370965e+06	5.40993748e-04	1.35371078e+06	6.56672045e-04	1.35371173e+06	6.56367359e-04
8387_pegase	2.74980929e+06	9.99884691e-03	2.77083829e+06	9.99896893e-03	2.77062704e+06	5.30460965e-02	2.77062704e+06	5.30460965e-02
9591_goc	1.02516095e+06	9.91659468e-04	1.06148769e+06	9.91997903e-04	1.06148806e+06	9.91795084e-04	1.06148807e+06	9.91788322e-04
9241_pegase	6.21775010e+06	4.18380648e-03	6.24208171e+06	4.18787958e-03	6.24207325e+06	3.76440386e-02	6.24207325e+06	3.76440386e-02
10480_goc	2.27696973e+06	1.09983709e-03	2.31442783e+06	1.09996886e-03	2.31442450e+06	1.67932256e-02	2.31442450e+06	1.67932256e-02
13659_pegase	8.92385389e+06	1.99904428e-03	8.94679835e+06	1.99980680e-03	8.94680070e+06	1.54477837e-02	8.94680070e+06	1.54477837e-02
19402_goc	1.93394723e+06	1.19983797e-03	1.97755237e+06	1.1999867e-03	1.97755235e+06	1.19986568e-03	1.97755235e+06	1.19986568e-03
24464_goc	2.58935630e+06	7.24722104e-04	2.62932336e+06	7.24944021e-04	2.62932439e+06	7.24724162e-04	2.62932439e+06	7.24724162e-04
30000_goc	1.11353160e+06	1.40161701e-03	1.14190983e+06	1.40292333e-03	1.14191122e+06	1.40225897e-03	1.14190714e+06	1.40184075e-03

- [10] F. Pacaud, S. Shin, M. Schanen, D. A. Maldonado, and M. Anitescu, “Accelerating condensed interior-point methods on simd/gpu architectures,” *Journal of Optimization Theory and Applications*, pp. 1–20, 2023.
- [11] D. Cole, S. Shin, F. Pacaud, V. M. Zavala, and M. Anitescu, “Exploiting gpu/simd architectures for solving linear-quadratic mpc problems,” in *2023 American Control Conference (ACC)*. IEEE, 2023, pp. 3995–4000.
- [12] C. G. Petra, N. Chiang, and J. Wang, “HiOp – User Guide,” Center for Applied Scientific Computing, Lawrence Livermore National Laboratory, Tech. Rep. LLNL-SM-743591, 2018.
- [13] N. Chiang, C. G. Petra, and V. M. Zavala, “Structured nonconvex optimization of large-scale energy systems using pips-nlp,” in *2014 Power Systems Computation Conference*. IEEE, 2014, pp. 1–7.
- [14] J. S. Rodriguez, R. B. Parker, C. D. Laird, B. L. Nicholson, J. D. Sirola, and M. L. Bynum, “Scalable parallel nonlinear optimization with pynumero and parapint,” *INFORMS Journal on Computing*, vol. 35, no. 2, pp. 509–517, 2023.
- [15] S. Shin, C. Coffrin, K. Sundar, and V. M. Zavala, “Graph-based modeling and decomposition of energy infrastructures,” *IFAC-PapersOnLine*, vol. 54, no. 3, pp. 693–698, 2021.
- [16] T. A. Davis and E. Palamadai Natarajan, “Algorithm 907: Klu, a direct sparse solver for circuit simulation problems,” *ACM Transactions on Mathematical Software (TOMS)*, vol. 37, no. 3, pp. 1–17, 2010.
- [17] “rosetta-opf,” <https://github.com/lanl-ansi/rosetta-opf>.
- [18] S. Babaeinejadsarookolae, A. Birchfield, R. D. Christie, C. Coffrin, C. DeMarco, R. Diao, M. Ferris, S. Fliscounakis, S. Greene, R. Huang *et al.*, “The power grid library for benchmarking ac optimal power flow algorithms,” *arXiv preprint arXiv:1908.02788*, 2019.
- [19] Y. Chen, T. A. Davis, W. W. Hager, and S. Rajamanickam, “Algorithm 887: Cholmod, supernodal sparse cholesky factorization and update/downdate,” *ACM Transactions on Mathematical Software (TOMS)*, vol. 35, no. 3, pp. 1–14, 2008.
- [20] L. Pineda, T. Fan, M. Monge, S. Venkataraman, P. Sodhi, R. T. Chen, J. Ortiz, D. DeTone, A. Wang, S. Anderson *et al.*, “Theseus: A library for differentiable nonlinear optimization,” *Advances in Neural Information Processing Systems*, vol. 35, pp. 3801–3818, 2022.

Government License: The submitted manuscript has been created by UChicago Argonne, LLC, Operator of Argonne National Laboratory (“Argonne”). Argonne, a U.S. Department of Energy Office of Science laboratory, is operated under Contract No. DE-AC02-06CH11357. The U.S. Government retains for itself, and others acting on its behalf, a paid-up nonexclusive, irrevocable worldwide license in said article to reproduce, prepare derivative works, distribute copies to the public, and perform publicly and display publicly, by or on behalf of the Government. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan. <http://energy.gov/downloads/doe-public-access-plan>.